

Introduction to Deep Learning

Convolutional Neural Network for Object Classification

John Luo

11/19/2018

Machine Learning?

- What does machine learning do?
- Tasks: **Classification, Regression, Transcription, Translation, Denoising**, etc.
- What is ML? Train a machine algorithm to do the above tasks.
 1. Start with **representations** of training data and model parameter.
 2. Training by **iterating** between mathematical **optimization** step and **evaluation** (minimizing training error).
 3. Apply trained algorithm onto test data (hoping the test error is small enough to generalize)

Some Optimization Concepts Revisited

- Bayesian statistics/**Baye's Rule**
- Maximum a posteriori (MAP) estimation
$$\boldsymbol{\theta}_{map} = \arg \max \log p(\boldsymbol{\theta}|\mathbf{x}) = \arg \max \log p(\mathbf{x}|\boldsymbol{\theta}) + \log p(\boldsymbol{\theta})$$
- Minimizing cost/lost function: such as the negative log probability density function above or any arbitrarily chosen function with regularization
- Some Regularizations:
 1. L2 regularization = Gaussian Prior
 2. L1 regularization = Laplace Prior
- **Gradient Descent Optimization**


Some ML Concepts

- **Supervised** Learning: training data is already tagged by human supervisors (classification and regression tasks).
- **Unsupervised** Learning: some underlying insight identified directly by machine learning programs (clustering, principal component analysis).
- We will focus on object classification using deep neural network as a case study on how deep learning works

| | <i>Supervised Learning</i> | <i>Unsupervised Learning</i> |
|-------------------|----------------------------------|------------------------------|
| <i>Discrete</i> | classification or categorization | clustering |
| <i>Continuous</i> | regression | dimensionality reduction |

Simple Neural Network Architecture

- On the right, NN of 1-hidden layer
- \mathbf{x}^i has 3 inputs (3 x1), \mathbf{W}^h is 3 x 4 matrix of weights connecting layers of neural nodes
- $\mathbf{W}^{iT} \mathbf{x}^i = \delta^h$
- Some activation function, g^h gives a pointwise operation: $g^h(\delta^h) = \mathbf{a}^h$
- To compute output layer, we have a matrix \mathbf{W}^o which is 4x2
- $\mathbf{W}^{oT} \mathbf{a}^h = \delta^o$
- Another activation function, g^o , a pointwise operation: $g^o(\delta^o) = \mathbf{y}^o$



Neural Network Basics

- Feedforward neural networks (often called multilayer perceptrons (MLPs)), each layer has linear transformation followed by an activation, 1st and 2nd layers:

$$\begin{aligned} \mathbf{h}^1 &= g^1(\mathbf{W}^{1T} \mathbf{x} + \mathbf{b}^1) \\ \mathbf{h}^2 &= g^2(\mathbf{W}^{2T} \mathbf{h}^1 + \mathbf{b}^2) \end{aligned}$$

- **W** represents **linear weights** assigned to each input, **b** represents **bias**, g is an **activation** function (can be nonlinear, see ReLU).
- Commonly used g are:
 - **Sigmoid** for binary classification (for Bernoulli output distributions):

$$g(z) = \frac{1}{1 + e^{-z}}$$

- **Softmax** for multiple classification (for multinoulli output distributions):

$$g(z)_i = \frac{\exp z_i}{\sum_j \exp z_j}$$


- **Rectified linear units (ReLU)**: $g(z) = \max\{0, z\}$, most widely used
- Tangent activation function: $g(z) = \tanh(z)$

Chain Rule & Backpropagation

- Chain rule is the key to **backpropagate error** and **update each layer's parameter W_i and b_i**
- The iterative update on W_i and b_i normally use a **gradient descent approach** with hyperparameters such as regularization (added to the gradients calculated) where needed and learning rate (step size of the gradient descent)
- Stochastic gradient descent (each sample assumed I.I.D, update parameters one sample at a time) is the most widely used, also full GD (use all samples) and batch GD (use a batch of sample to update)
- Regularization: L1, L2 or arbitrary
- Also check out enhanced gradient descent algorithms such as Nestorov momentum and RMSprop (for example, stochastic GD with RMSprop)
- Optimal hyperparameters can be determined using **grid search** or **random search** (a time-consuming guessing game)


Convolutional Neural Network

- Convolution operation is very useful in processing 2D images because it is **translation invariant** and **computationally inexpensive** compared to a matrix operation
- Convolution is a linear transformation, replacing $\mathbf{W}x + \mathbf{b}$ in regular feedforward neural networks (often use “same” size option to specify convolution so output size = input size)
- Nonlinearity step such as **Rectified Linear Unit** still used
- **New building block: Pooling** such as max pooling:




A Simple CNN Architecture

1. RGB image $H \times W \times 3$ ($224 \times 224 \times 3$), we want to know what is in the input image
2. First **convolution** filter stack of $3 \times K \times K \times F$: a stack of F filters each detect a feature (for example, detecting horizontal and vertical edges will require a stack of $3 \times K \times K \times 2$).
3. Then **ReLU** is applied pointwise.
4. Second **convolution** filter stack is $64 \times K \times K \times 64$, third one is $64 \times K \times K \times 128$, then **ReLU** again.
5. followed by **max pooling of 2×2** to reduce $224 \times 224 \times 128$ to $112 \times 112 \times 128$, repeat Steps (2 to 5)
6. Repeat from Step 2 to 5 with different matrices' size
7. After all Convo+ReLU layers, **stack the $7 \times 7 \times 512$ to $1 \times 1 \times 4096$, called "Flatten"**
8. In the end, a **softmax** operation can be used to do **classification** of objects, here identifying possibly 1000 objects



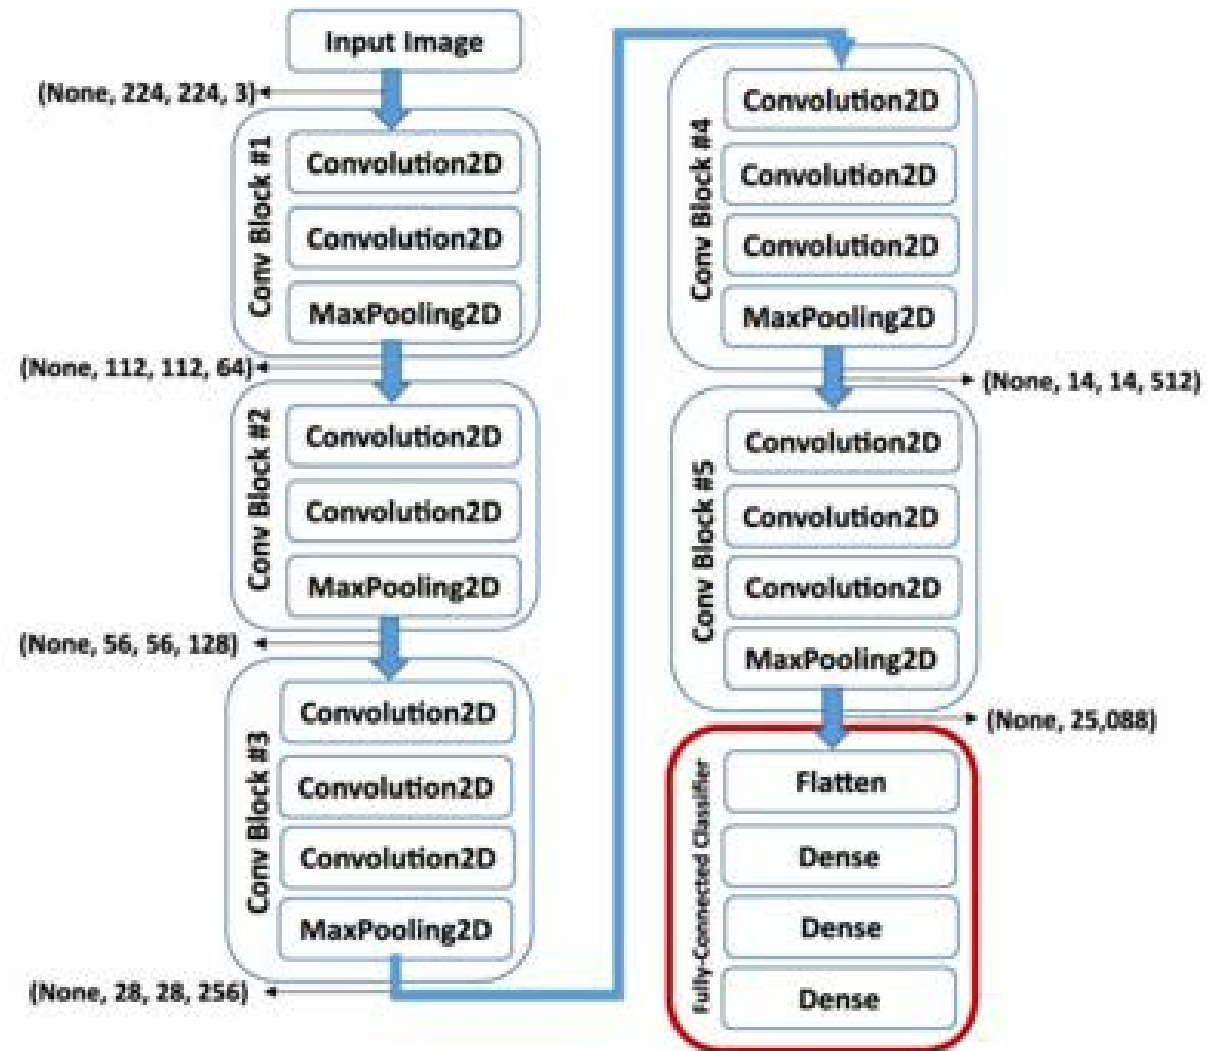
How to Interpret this?

- The first convolution $3 \times K \times K \times \mathbf{F}$ extract **\mathbf{F} number of features**, these features are simple features such as lines and edges.
- The output of the convolution operation indicates where the feature is found, max pooling retain that info while keeping the data size small (down sampling).
- Here the 1st convolution $3 \times K \times K \times \mathbf{64}$ extract **64 simple features** (shapes) from the RGB image
- The 3rd convolution $64 \times K \times K \times \mathbf{128}$ extracts more complicated features (128 of them) on top of the simple ones
- The deepest convolution stack ($512 \times K \times K \times \mathbf{512}$) **extracts very sophisticated feature representation** such as a car, a tree
- The **fully connected + ReLU layer** maps the stacked output of convolutional layers to an identification matrix: logits indicating if a car or a tree is present in the input




Some Widely Used CNN Architectures

- The training of CNN uses chain rule as basis again
- Basically the deeper the better up to some extent
- All are combo of Conv and ReLU and Pooling with the last few layers as feedforward neural layers
- **VGG16 architecture on the right**
- Can find pre-trained model online




ResNet

- Deeper is not always better, for example identity operation $f(x) = x$ is notoriously difficult to replicate
- Introducing residual block




34-layer residual



34-layer plain



VGG-19



Practical Implementation/Software Tools

- Python with **sklearn**, **Theano**, **TensorFlow**(google) libraries. **Keras**, written using TensorFlow is a high-level neural network library (very simple to set up complicated neural networks with)
- C++ with **Caffe** library
- MATLAB with **MatConvNet** library
- All the above can achieve **GPU integration** to take advantage of its faster computational speed
- Pre-trained weights of an image classifier (for example a pretrained VGG16 model) can be downloaded from ImageNet

Onto Computer Vision

- We also want to identify the location of the object that we identified.
- Modify the final layers of fully connected neural networks (only used to determine if an object is there) and incorporate linear regression to get center coordinates and height and width
- One of the state-of-art algorithm is called **Single-Shot MultiBox Detector (SSD)** which adds different box predictors along with the original class predictors at different scales/default boxes to detect and locate objects

